

# Data Structure

Michael  
2017

## Data Structure

VPP16.09

```
typedef struct _vlib_node_registration
{
    vlib_node_function_t *function;
    char *name;
    char *sibling_of;
    u32 index;
    vlib_node_type_t type;
    char **error_strings;
    format_function_t *format_buffer;
    unformat_function_t *unformat_buffer;
    format_function_t *format_trace;
    unformat_function_t *unformat_trace;
    u8 *(*validate_frame)(struct vlib_main_t * vm,
                          struct vlib_node_runtime_t *,
                          struct vlib_frame_t * f);

    void *runtime_data;
    u16 process_log2_n_stack_bytes;
    u8 runtime_data_bytes;
    u8 state;
    u16 flags;
    u16 scalar_size, vector_size;
    u16 n_errors;
    u16 n_next_nodes;
    struct _vlib_node_registration *next_registration;
    char *next_nodes[];
} vlib_node_registration_t;

typedef struct vlib_node_t
{
    vlib_node_function_t *function;
    u8 *name;
    u32 name_eolog_string;
    vlib_node_stats_t stats_total;
    vlib_node_stats_t stats_last_clear;
    vlib_node_type_t type;
    u32 index;
    u32 runtime_index;
    void *runtime_data;
    u16 flags;
    u8 state;
    u8 runtime_data_bytes;
    u16 n_errors;
    u16 scalar_size, vector_size;
    u32 error_heap_handle;
    u32 error_heap_index;
    char **error_strings;
    char **next_node_names;
    u32 *next_nodes;
    char *sibling_of;
    uword *sibling_bitmap;
    u64 *n_vectors_by_next_node;
    uword *next_slot_by_node;
    uword *prev_node_bitmap;
    u32 owner_node_index, owner_next_index;
    format_function_t *format_buffer;
    unformat_function_t *unformat_buffer;
    format_function_t *format_trace;

    u8 *(*validate_frame)(struct vlib_main_t * vm,
                          struct vlib_node_runtime_t *,
                          struct vlib_frame_t * f);

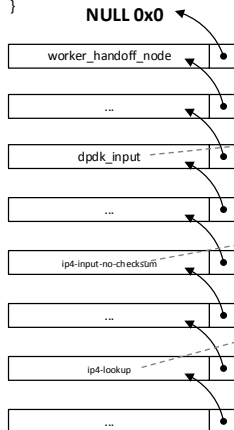
    u8 *state_string;
} vlib_node_t;
```

```
typedef struct
{
    vlib_node_t **nodes;
    uword *node_by_name;
    u32 flags;
    vlib_node_runtime_t *nodes_by_type[VLIB_N_NODE_TYPE];
    u32 *pending_interrupt_node_runtime_indices;
    u32 polling_threshold_vector_length;
    u32 interrupt_threshold_vector_length;
    vlib_next_frame_t *next_frames;
    vlib_pending_frame_t *pending_frames;
    timing_wheel_t timing_wheel;
    vlib_signal_timed_event_data_t *signal_timed_event_data_pool;
    u32 *data_from_advancing_timing_wheel;
    u64 cpu_time_next_process_ready;
    vlib_process_t **processes;
    u32 current_process_index;
    vlib_pending_frame_t *suspended_process_frames;
    void **recycled_event_data_vectors;
    u32 input_node_counts_by_state[VLIB_N_NODE_STATE];
    uword *frame_size_hash;
    vlib_frame_size_t *frame_sizes;
    f64 time_last_runtime_stats_clear;
    vlib_node_registration_t *node_registrations;
} vlib_node_main_t;
```

```
typedef struct vlib_node_runtime_t
{
    vlib_node_function_t *function;
    vlib_error_t *errors;
    u32 clocks_since_last_overflow;
    u32 max_clock;
    u32 max_clock_n;
    u32 calls_since_last_overflow;
    u32 vectors_since_last_overflow;
    u32 next_frame_index;
    u32 node_index;
    u32 input_main_loops_per_call;
    u32 main_loop_count_last_dispatch;
    u32 main_loop_vector_stats[2];
    u16 flags;
    u16 state;
    u16 n_next_nodes;
    u16 cached_next_index;
    u16 cpu_index;
    uword runtime_data[128]
} vlib_node_runtime_t;
```

- 1 \* sizeof (vlib\_node\_function\_t \*)  
 - 1 \* sizeof (vlib\_error\_t \*)  
 - 11 \* sizeof (u32)  
 - 5 \* sizeof (u16) / sizeof (uword));

```
VLIB_REGISTER_NODE
{
    vlib_main_t * vm = vlib_get_main();
    x.next_registration = vm->node_main.node_registrations;
    vm->node_main.node_registrations = &x;
}
```

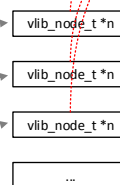


vm->node\_main.node\_registrations  
 Copyright© 2017. All rights reserved.

```
vlib_register_all_static_nodes (vlib_main_t * vm)
{
    vlib_node_registration_t *r;

    r = vm->node_main.node_registrations;
    while (r)
    {
        register_node (vm, r);
        r = r->next_registration;
    }
}
```

```
heap = clib_per_cpu_mheaps[cpu];
n = clib_mem_alloc_no_fail (sizeof (n[0]));
```



```
n->index = vec_len (nm->nodes);
vec_add1 (nm->nodes, n);
```

```
hash_set (nm->node_by_name, n->name, n->index);
```

## VPP16.09

- 

On heap [Y]

$\rightarrow$  nm->nodes\_by\_type[n->type]

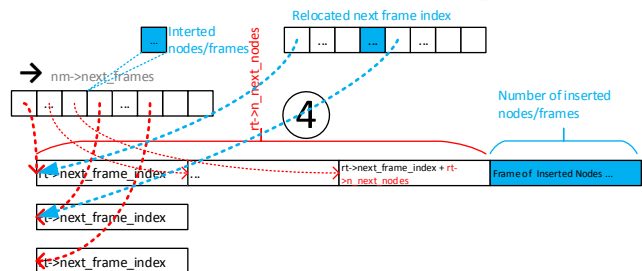
```
vec_add2_aligned (nm->nodes_by_type[n->type], rt, 1,
/* align */ CUB_CACHE_LINE_BYTES);
n->runtime_index = rt - nm->nodes_by_type[n->type];
```

vlib\_node\_runtime\_t \*rt

vlib\_node\_runtime\_t \*rt

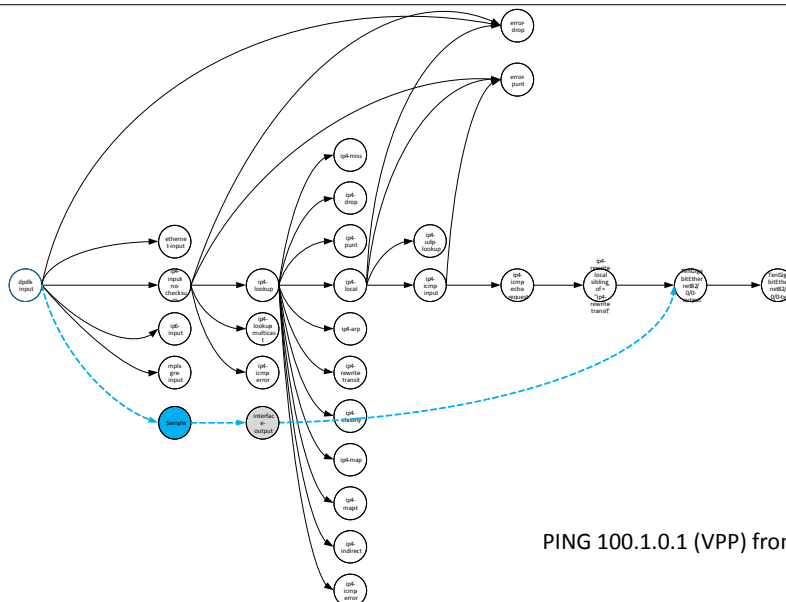
vlib\_node\_runtime\_t \*rt

...



5

## VPP16.09



PING 100.1.0.1 (VPP) from 100.1.0.2(Host)

6

## Register Node and Runtime Graph(Cont')

VPP16.09

```

/* *INDENT-OFF* */
VLIB_REGISTER_NODE (dpdk_input_node) = {
    .function = dpdk_input,
    .type = VLIB_NODE_TYPE_INPUT,
    .name = "dpdk-input",

    /* Will be enabled if/when hardware is detected. */
    .state = VLIB_NODE_STATE_DISABLED,

    .format_buffer = format_ethernet_header_with_length,
    .format_trace = format_dpdk_rx_dma_trace,

    .n_errors = DPDK_N_ERROR,
    .error_strings = dpdk_error_strings,

    .n_next_nodes = DPDK_RX_N_NEXT,
    .next_nodes = {
        [DPDK_RX_NEXT_DROP] = "error-drop",
        [DPDK_RX_NEXT_ETHERNET_INPUT] = "ethernet-input",
        [DPDK_RX_NEXT_IP4_INPUT] = "ip4-input-no-checksum",
        [DPDK_RX_NEXT_IP6_INPUT] = "ip6-input",
        [DPDK_RX_NEXT_MPLS_INPUT] = "mpls-gre-input",
    },
};

VLIB_REGISTER_NODE
(ip4_input_no_checksum_node,static) = {
    .function = ip4_input_no_checksum,
    .name = "ip4-input-no-checksum",
    .vector_size = sizeof (u32),

    .n_next_nodes = IP4_INPUT_N_NEXT,
    .next_nodes = {
        [IP4_INPUT_NEXT_DROP] = "error-drop",
        [IP4_INPUT_NEXT_PUNT] = "error-punt",
        [IP4_INPUT_NEXT_LOOKUP] = "ip4-lookup",
        [IP4_INPUT_NEXT_LOOKUP_MULTICAST] = "ip4-lookup-
multicast",
        [IP4_INPUT_NEXT_ICMP_ERROR] = "ip4-icmp-error",
    },

    .format_buffer = format_ip4_header,
    .format_trace = format_ip4_input_trace,
};

VLIB_REGISTER_NODE (ip4_lookup_node) = {
    .function = ip4_lookup,
    .name = "ip4-lookup",
    .vector_size = sizeof (u32),

    .format_trace = format_ip4_lookup_trace,

    .n_next_nodes = IP4_LOOKUP_N_NEXT,
    .next_nodes = IP4_LOOKUP_NEXT_NODES,
};

```

Copyright© 2017. All rights reserved.

Michael

7

## Register Node and Runtime Graph(Cont')

VPP16.09

```

VLIB_REGISTER_NODE (ip4_local_node,static) = {
    .function = ip4_local,
    .name = "ip4-local",
    .vector_size = sizeof (u32),

    .format_trace = format_ip4_forward_next_trace,

    .n_next_nodes = IP_LOCAL_N_NEXT,
    .next_nodes = {
        [IP_LOCAL_NEXT_DROP] = "error-drop",
        [IP_LOCAL_NEXT_PUNT] = "error-punt",
        [IP_LOCAL_NEXT_UDP_LOOKUP] = "ip4-udp-lookup",
        [IP_LOCAL_NEXT_ICMP] = "ip4-icmp-input",
    },
};

VLIB_REGISTER_NODE (ip4_icmp_input_node,static) = {
    .function = ip4_icmp_input,
    .name = "ip4-icmp-input",

    .vector_size = sizeof (u32),

    .format_trace = format_icmp_input_trace,

    .n_errors = ARRAY_LEN (icmp_error_strings),
    .error_strings = icmp_error_strings,

    .n_next_nodes = 1,
    .next_nodes = {
        [ICMP_INPUT_NEXT_ERROR] = "error-punt",
    },
};

VLIB_REGISTER_NODE
(ip4_icmp_echo_request_node,static) = {
    .function = ip4_icmp_echo_request,
    .name = "ip4-icmp-echo-request",

    .vector_size = sizeof (u32),

    .format_trace = format_icmp_input_trace,

    .n_next_nodes = 1,
    .next_nodes = {
        [0] = "ip4-rewrite-local",
    },
};

VLIB_REGISTER_NODE (ip4_rewrite_local_node) = {
    .function = ip4_rewrite_local,
    .name = "ip4-rewrite-local",
    .vector_size = sizeof (u32),

    .sibling_of = "ip4-rewrite-transit",

    .format_trace = format_ip4_rewrite_trace,

    .n_next_nodes = 0,
};

CLIB_MULTIARCH_SELECT_FN
(vnet_interface_output_node_no_flatten);

VNET_DEVICE_CLASS (dpdk_device_class) = {
    .name = "dpdk",
    .tx_function = dpdk_interface_tx,
    .tx_function_n_errors = DPDK_TX_FUNC_N_ERROR,
    .tx_function_error_strings = dpdk_tx_func_error_strings,
    .format_device_name = format_dpdk_device_name,
    .format_device = format_dpdk_device,
    .format_tx_trace = format_dpdk_tx_dma_trace,
    .clear_counters = dpdk_clear_hw_interface_counters,
    .admin_up_down_function =
dpdk_interface_admin_up_down,
    .subif_add_del_function = dpdk_subif_add_del_function,
    .rx_redirect_to_node = dpdk_set_interface_next_node,
    .no_flatten_output_chains = 1,
    .name_renumber = dpdk_device_renumber,
};

```

Copyright© 2017. All rights reserved.

Michael

8