

# Vlib Buffer

Michael  
2017

## The abstract of Dave's mail

VPP16.09

Vlib\_buffer\_t's are embedded in private headroom space, tangent to the packet DMA target [which starts at (struct mbuf \*)mb->data\_off].

Here's how we create the mempools:

```
rmf = rte_pktmbuf_pool_create ((char *) pool_name, /* pool name */
    num_mbufs, /* number of mbufs */
    512, /* cache size */
    VLIB_BUFFER_HDR_SIZE, /* priv size */
    VLIB_BUFFER_PRE_DATA_SIZE + VLIB_BUFFER_DATA_SIZE, /* dataroom size */
    socket_id); /* cpu socket */
```

As Chris wrote, b0->data is an array, not a pointer:

```
u8 data[0]; /**< Packet data. Hardware DMA here */
```


Q: How to manipulate packet payload in VPP?

See vlib\_buffer\_advance(...). Depending on what you want to do - which you didn't share - that may be all you need.

In any event, these buffer metadata fields are involved:

current\_data, current\_length, total\_length\_not\_including\_first\_buffer, next\_buffer, flags

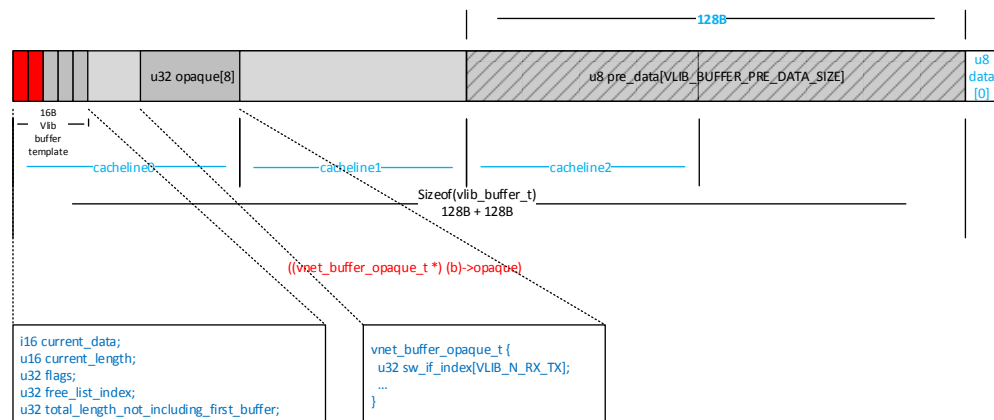
Take a careful look at .../src/vlib/{buffer.h,buffer\_funcs.h} to see usage patterns.




```

rmp = rte_pktmbuf_pool_create ((char *) pool_name, /* pool name */
                               num_mbufs, /* number of mbufs */
                               512, /* cache size */
                               VLIB_BUFFER_HDR_SIZE, /* priv size */
                               VLIB_BUFFER_PRE_DATA_SIZE + VLIB_BUFFER_DATA_SIZE, /* dataroom size */
                               socket_id); /* cpu socket */

```





```
b0->error = node->errors[error0];
```

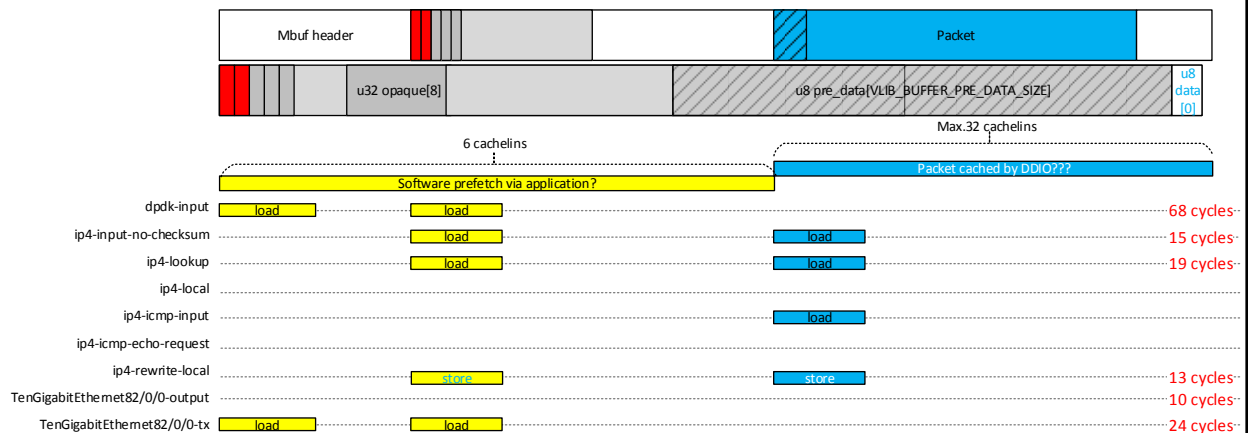
```
l3_offset0 = ((next0 == DPDK_RX_NEXT_IP4_INPUT ||
               next0 == DPDK_RX_NEXT_IP6_INPUT ||
               next0 == DPDK_RX_NEXT_MPLS_INPUT) ?
              sizeof(ethernet_header_t) : 0);
```

```
b0->current_data = l3_offset0;
/* Some drivers like fm10k receive frames with
   mb->data_off > RTE_PKTMBUF_HEADROOM */
b0->current_data += mb->data_off - RTE_PKTMBUF_HEADROOM;
b0->current_length = mb->data_len - l3_offset0;
```

```
b0->flags = buffer_flags_template;
```

```
if (VMWARE_LENGTH_BUG_WORKAROUND)
    b0->current_length -= 4;
```

```
vnet_buffer(b0)->sw_if_index[VLIB_RX] = xd->vlib_sw_if_index;
vnet_buffer(b0)->sw_if_index[VLIB_TX] = (u32) ~ 0;
n_rx_bytes += mb->pkt_len;
```



Level	size	Latency	IndexBits	Sets	Associativity	LineSize
L1Data	32KiB	4	6	64	8	64B
L1Instructions	32KiB	4	6	64	8	64B
L2	256KiB	7	9	512	8	64B
LLC	≥3MiB	26-31	11	≥4096	12-20	64B

Typical cache parameters in Intel processors