

Trace

Michael
2017

Data Structure

VPP16.09

```
typedef struct
{
    /* CPU time stamp trace was made. */
    u64 time;

    /* Node which generated this trace. */
    u32 node_index;

    /* Number of data words in this trace. */
    u32 n_data;

    /* Trace data follows. */
    u8 data[0];
} vlib_trace_header_t;
```

Flags

- VLIB_NODE_FLAG_TRACE
- VLIB_BUFFER_IS_TRACED

DBGvpp# trace filter include error-drop 1

cli_filter_trace()

DBGvpp# trace add dpdk-input 100

cli_add_trace_buffer()

typedef struct

{

u32 count;

u32 limit;

} vlib_trace_node_t;

DBGvpp# show trace

cli_show_trace_buffer()

- trace_apply_filter() <--- filter applied here

-- filter_accept()

Trace functions:**Vlib_trace_...()****Filter:****trace_apply_filter()**

...

Add Trace

always_inline void *

vlib_add_trace (vlib_main_t * vm,vlib_node_runtime_t * r, **vlib_buffer_t * b**, u32 n_data_bytes)

{

vlib_trace_main_t *tm = &vm->trace_main;

vlib_trace_header_t *h;

u32 n_data_words;

vlib_validate_trace (tm, b);

n_data_bytes = round_pow2 (n_data_bytes, sizeof (h[0]));

n_data_words = n_data_bytes / sizeof (h[0]);

vec_add2_aligned (**tm->trace_buffer_pool[b->trace_index]**, h,
1 + n_data_words, sizeof (h[0]));

→tm->trace_buffer_pool[]

h->time = vm->cpu_time_last_node_dispatch;

h->n_data = n_data_words;

h->node_index = r->node_index;

return h->data;

}

```
vlib_trace_buffer(vlib_main_t * vm,
                  vlib_node_runtime_t * r,
                  u32 next_index, vlib_buffer_t * b, int follow_chain)
{
    ...
    pool_get (tm->trace_buffer_pool, h);

    do
    {
        b->flags |= VLIB_BUFFER_IS_TRACED;
        b->trace_index = h - tm->trace_buffer_pool;
    }
    while (follow_chain && (b = vlib_get_next_buffer (vm, b)));
}
```

```
t0 = vlib_add_trace (vm, node, b0, sizeof (t0[0]));
t0->adj_index = vnet_buffer (b0)->ip.adj_index[which_adj_index];
t0->flow_hash = vnet_buffer (b0)->ip.flow_hash;
t0->fib_index = (vnet_buffer(b0)->sw_if_index[VLIB_TX] != (u32)~0) ?
    vnet_buffer(b0)->sw_if_index[VLIB_TX] :
    vec_elt (im->fib_index_by_sw_if_index,
            vnet_buffer(b0)->sw_if_index[VLIB_RX]);

clib_memcpy (t0->packet_data,
            vlib_buffer_get_current (b0),
            sizeof (t0->packet_data));
```

```

typedef struct
{
    u32 buffer_index;
    u16 device_index;
    u16 queue_index;
    struct rte_mbuf mb;
    vlib_buffer_t buffer;
    u8 data[256];
} dpdk_rx_dma_trace_t;

15:12:46:260666: dpdk-input
TenGigabitEthernet82/0/0 rx queue 0
buffer 0x1b390: current data 14, length 84, free-list 0, totlen-nifb
0, trace 0x1
PKT MBUF: port 0, nb_segs 1, pkt_len 98
buf_len 2176, data_len 98, ol_flags 0x0, data_off 128, phys_addr
0x7b587540
packet_type 0x10
Packet Types
RTE_PTYPE_L3_IPV4 (0x0010) IPv4 packet without extension headers
IP4: 90:e2:ba:84:1c:3a -> 90:e2:ba:84:45:b2
ICMP: 100.1.0.2 -> 100.1.0.1
tos 0x00, ttl 64, length 84, checksum 0x43d4
fragment id 0x2ed0, flags DONT_FRAGMENT
ICMP echo_request checksum 0x500c

typedef struct {
    u8 packet_data[64];
} ip4_input_trace_t;

15:12:46:260675: ip4-input-no-checksum
ICMP: 100.1.0.2 -> 100.1.0.1
tos 0x00, ttl 64, length 84, checksum 0x43d4
fragment id 0x2ed0, flags DONT_FRAGMENT
ICMP echo_request checksum 0x500c

typedef struct {
    u32 adj_index;
    u32 flow_hash;
    u32 fib_index;
    u8 packet_data[64 -
1*sizeof(u32)];
} ip4_forward_next_trace_t;

15:12:46:260683: ip4-lookup
fib 0 adj-idx 4 : 100.1.0.1/24 flow hash: 0x00000000
ICMP: 100.1.0.2 -> 100.1.0.1
tos 0x00, ttl 64, length 84, checksum 0x43d4
fragment id 0x2ed0, flags DONT_FRAGMENT
ICMP echo_request checksum 0x500c

```

Copyright© 2017. All rights reserved.

Michael

7

How to Trace Packets Which VPP Creates

Folks,

If one wants/needs to trace packets which vpp creates from whole cloth, here's how to do it. Typically, it's an input-node function.

```

/* Top of node dispatch function */
u32 n_trace = vlib_get_trace_count (vm, node);

/* per-buffer */
VLIB_BUFFER_TRACE_TRAJECTORY_INIT (b0);

/* If actually tracing generated packets... */
if (PREDICT_FALSE (n_trace > 0))
{
    session_queue_trace_t *t0;
    vlib_trace_buffer (vm, node, next_index, b0,
        1 /* follow_chain */ );
    vlib_set_trace_count (vm, node, --n_trace);
    t0 = vlib_add_trace (vm, node, b0, sizeof (*t0));

    <fill in rest of trace>
}

```

In case of interior graph nodes, packets came from somewhere, and should be traced from where they entered the graph. Otherwise, information is lost.

Thanks... Dave

Copyright© 2017. All rights reserved.

Michael

8